

# LEGES

parallel R  
in the cloud  
two lines of code  
  
no kidding!

why...

so i've got this problem...

reinsurance simulations  
updated frequently for one month

on my laptop...

each sim takes ~ 1 min

10k sims \* 1 min = ~ 7 days

no need for full map/reduce  
embarrassingly parallel

# W U G E S

you've seen "word count"  
demos...

segue has nothing to do  
with that

big cpu, not big data

# my options...

SEQUENCE

make the code faster

build a cluster

type

snow

mpi

hadoop

location

self hosted

amazon web services

ec2

emr

rackspace

```
graph LR; subgraph Options; direction TB; type["type"]; type --- snow["snow"]; type --- mpi["mpi"]; type --- hadoop["hadoop"]; location["location"]; location --- self_hosted["self hosted"]; location --- aws["amazon web services"]; location --- ec2["ec2"]; location --- emr["emr"]; location --- rackspace["rackspace"]; end; hadoop --> costs["lowest startup costs"]; emr --> costs;
```

lowest startup costs

# SEGUE

what did my mind's eye see?

# SEQUE



"simply irresistible" ~8000 hp

# SEQUE

syntax...

```
require(segue)
```

```
myCluster <- createCluster()
```

contratulations. we've built a hadoop  
cluster!

more syntax...

parallel apply() on lists:

base R:

`lapply( X, FUN, ... )`


segue:

`emrlapply( clusterObject, X, FUN, ... )`



example...

stochastic pi simulation (again)



# example...

```
estimatePi <- function( seed ){  
  set.seed(seed)  
  numDraws <- 1000000  
  r <- .5  
  x <- runif(numDraws, min=-r, max=r)  
  y <- runif(numDraws, min=-r, max=r)  
  inCircle <- ifelse( (x^2 + y^2)^.5 < r , 1, 0)  
  return(sum(inCircle) / length(inCircle) * 4)  
}
```

```
seedList <- as.list(1:1000)
```

```
require(segue)
```

```
myCluster <- createCluster(20)
```

```
myEstimates <- emrapply( myCluster, seedList, estimatePi )
```

```
stopCluster(myCluster)
```

```
myPi <- Reduce(sum, myEstimates) / length(myEstimates)
```

```
format(myPi, digits=10)
```

<https://gist.github.com/764370>

SEGUE

# howzit work?

## SEQUENCE

### createCluster()

cluster object:  
list of parameters

temp dirs:  
local  
S3 for EMR

bootstrap:  
update R  
update packages

~ 10-15 minutes

# howzit work?

## emrapply()

# SEQUENCE

list is serialized to CSV and uploaded to S3 – streaming input file

function, arguments, r objects, etc are saved & uploaded


EMR copies files to nodes – mapper.R picks them up

CSV is input to mapper.R  
applies function to each list element

output is serialized into emr part-xxxxx files on s3

part files are downloaded to R and deserialized

deserialized results are reordered and put into a list object



# SEQUE

```
createCluster( numInstances=2,  
               cranPackages,  
               filesOnNodes,  
               rObjectsOnNodes,  
               enableDebugging=FALSE,  
               instancesPerNode,  
               masterInstanceType="m1.small",  
               slaveInstanceType="m1.small",  
               location="us-east-1a",  
               ec2KeyName,  
               copy.image=FALSE,  
               otherBootstrapActions,  
               sourcePackagesToInstall)
```

numInstances  
cranPackages  
filesOnNodes  
rObjectsOnNodes  
enableDebugging  
instancesPerNode  
masterInstanceType  
slaveInstanceType  
location  
ec2KeyName  
copy.image  
otherBootstrapActions  
sourcePackagesToInstall

number of ec2 machines to fire up  
cran packages to load on each cluster node  
files to be loaded on each node  
R objects to put on the worker nodes  
start emr debugging  
number of R instances per node  
ec2 instance type for the master node  
ec2 instance type for the slave nodes  
ec2 location name for the cluster  
ec2 key used for logging into the main node  
copy the entire local environment to the nodes?  
other bootstrap actions to run  
R source packages to be installed on each node

# when to use segue...

W  
U  
G  
E  
S

embarrassingly parallel

cpu bound

apply on lists with many items

object size: to / from s3 roundtrip

each job has a fixed & marginal cost

# SEGUE

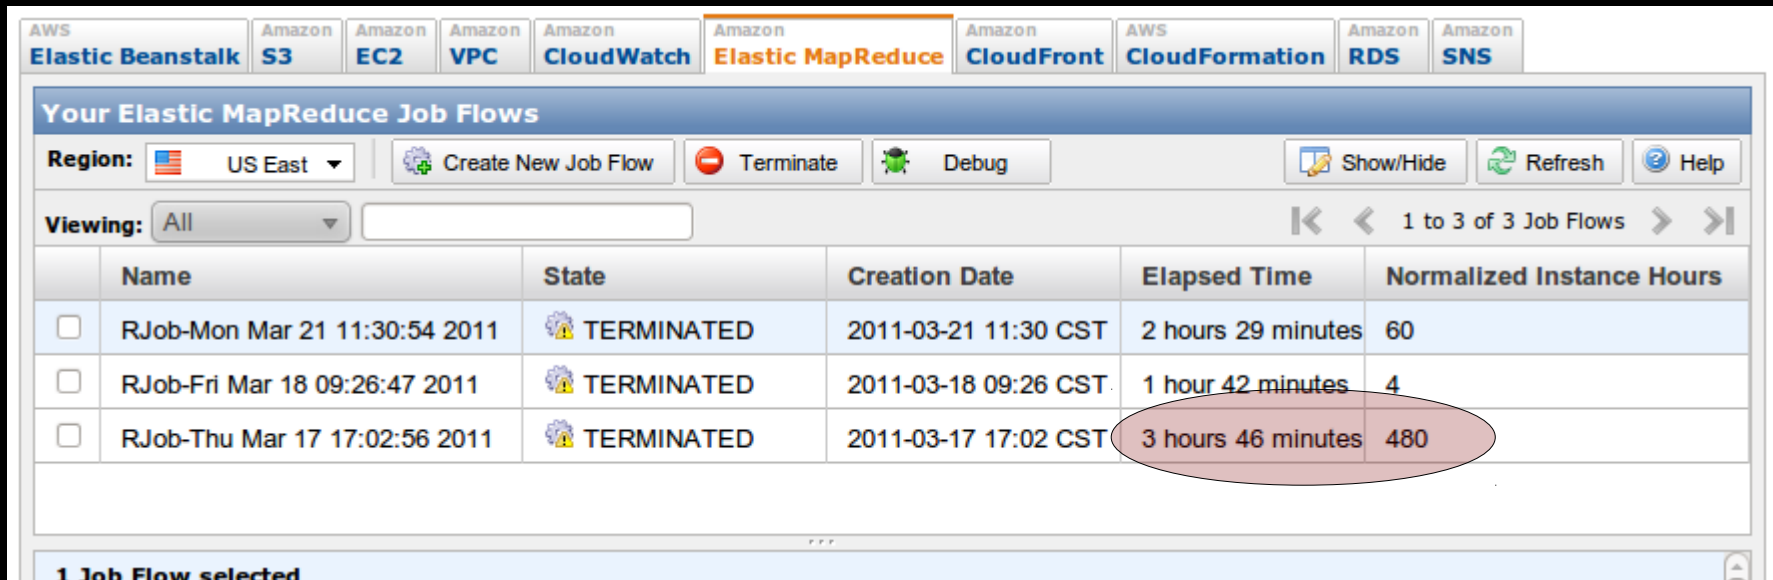
downside of segue...  
embarrassingly parallel failure



SE  
GUE

reasons daddy drinks...  
a.k.a things vendors never say

keep one eye on aws dashboard



Region: US East

Viewing: All

|                          | Name                          | State      | Creation Date        | Elapsed Time       | Normalized Instance Hours |
|--------------------------|-------------------------------|------------|----------------------|--------------------|---------------------------|
| <input type="checkbox"/> | RJob-Mon Mar 21 11:30:54 2011 | TERMINATED | 2011-03-21 11:30 CST | 2 hours 29 minutes | 60                        |
| <input type="checkbox"/> | RJob-Fri Mar 18 09:26:47 2011 | TERMINATED | 2011-03-18 09:26 CST | 1 hour 42 minutes  | 4                         |
| <input type="checkbox"/> | RJob-Thu Mar 17 17:02:56 2011 | TERMINATED | 2011-03-17 17:02 CST | 3 hours 46 minutes | 480                       |

1 Job Flow selected

united nations considers debugging of  
segue jobs "torture" under geneva  
convention



more reasons daddy drinks...

SEGUE

if you use segue you will see:  
unreproducible errors  
clusters that never start  
temp buckets in your s3 acct  
clusters left running  
i/o that takes longer than calcs

but... i've never had a "wrong"  
answer

# SEQUE

immediate segue future...

maintenance issues:  
R releases change  
emr changes

vendor lock-in to amazon

whirr as solution?

foreach %dopar% backend?

imagine the future...

R objects backed by clusters  
`as.hdfs.data.frame(data)`

operations converted to map reduce  
 jobs transparently

abstractions...

# W U G E S

segue project page

<http://code.google.com/p/segue/>

google groups


<http://groups.google.com/group/segue-r>

see also...

rhipe – program m/r in R

<http://www.stat.purdue.edu/~sguha/rhipe/>

## R Cookbook...



### R Cookbook

With more than 200 practical recipes, this book helps you perform data analysis with R quickly and efficiently. The R language provides everything you need to do statistical work, but its structure can be difficult to master. This collection of concise, task-oriented recipes makes you productive with R immediately, with solutions ranging from basic tasks to input and output, general statistics, graphics, and linear regression.

Each recipe addresses a specific problem, with a discussion that explains the solution and offers insight into how it works. If you're a beginner, *R Cookbook* will help get you started. If you're an experienced data programmer, it will jog your memory and expand your horizons. You'll get the job done faster and learn more about R in the process.

- Create vectors, handle variables, and perform other basic functions
- Input and output data
- Tackle data structures such as matrices, lists, factors, and data frames
- Work with probability, probability distributions, and random variables
- Calculate statistics and confidence intervals, and perform statistical tests
- Create a variety of graphic displays
- Build statistical models with linear regressions and analysis of variance (ANOVA)
- Explore advanced statistical techniques such as finding clusters in your data

*“Wonderfully readable, R Cookbook serves not only as a solutions manual of sorts, but as a truly enjoyable way to explore the R language—one practical example at a time.”*

—Jeffrey Ryan  
Software consultant and  
R package author

*“With 95% confidence, I fail to reject that R Cookbook is the best book for learning and using the important stats functions in R.”*

—JD Long  
R Blogger at  
CerebralMastication.com

Paul Teetor is a quantitative developer with Masters degrees in statistics and computer science. He specializes in analytics and software engineering for investment management, securities trading, and risk management..